

Name	Sreeshma.S
Register Number	720419104024
Team Size	4
Team ID	PNT2022TMID43507



ASSIGNMENT-II

```
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import numpy as np
sns.set_style('darkgrid')
sns.set(font_scale=1.3)
```

```
df=pd.read_csv("/content/drive/MyDrive/IBM/Assignment - 2
/Churn_Modelling.csv")
df.head()
```

```
RowNumber CustomerId Surname CreditScore Geography Gender Age
```

```
\
```

```
0          1  15634602 Hargrave      619  France Female  42
1          2  15647311   Hill      608  Spain Female  41
2          3  15619304   Onio      502  France Female  42
3          4  15701354   Boni      699  France Female  39
4          5  15737888 Mitchell     850  Spain Female  43
```

```
Tenure  Balance  NumOfProducts  HasCrCard  IsActiveMember \
1         1         1
1         1  83807.86          1         0         1
2         8 159660.80          3         1         0
3         1    0.00          2         0         0
4         2 125510.82          1         1         1
```

```
EstimatedSalary Exited
101348.88         1
```

```

1          112542.58    0
2          113931.57    1
3          93826.63    0 4          79084.10    0

```

```
df.drop(["RowNumber","CustomerId","Surname"],axis=1,inplace=True) df.info()
```

```
<class 'pandas.core.frame.DataFrame'> RangeIndex: 10000
entries, 0 to 9999
```

Data columns (total 11 columns):

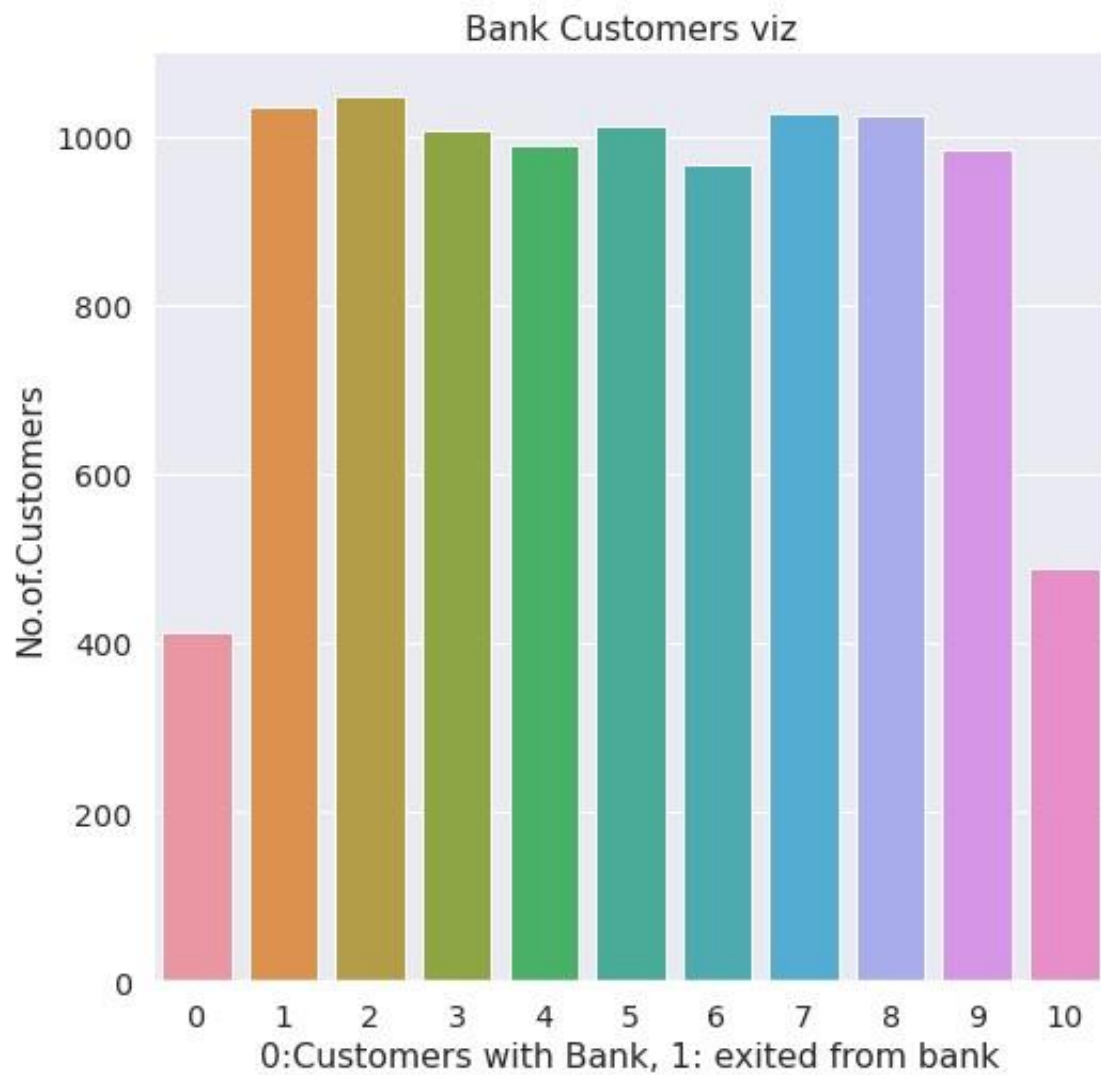
```

# Column      Non-Null Count  Dtype
---  -
0  CreditScore    10000 non-null int64  1  Geography
10000 non-null object  2  Gender      10000 non-null
object  3  Age          10000 non-null int64
4  Tenure         10000 non-null int64  5  Balance
10000 non-null float64  6  NumOfProducts 10000 non-
null int64
7  HasCrCard      10000 non-null int64  8
IsActiveMember  10000 non-null int64  9
EstimatedSalary 10000 non-null float64 10  Exited
10000 non-null int64 dtypes: float64(2), int64(7),
object(2) memory usage: 859.5+ KB

```


```
#Perform Univariate Analysis plt.figure(figsize=(8,8))
sns.countplot(x='Tenure',data=df)
```

```
plt.xlabel('0:Customers with Bank, 1: exited from bank')
plt.ylabel('No.of.Customers') plt.title("Bank Customers viz") plt.show()
```




```
#Perform Univariate Analysis plt.figure(figsize=(8,8))  
sns.kdeplot(x=df['Balance'])
```

<matplotlib.axes._subplots.AxesSubplot at 0x7fa0c03906d0>



```
#Perform Bivariate Analysis plt.scatter(df.Age,df.Balance)
```

```
<matplotlib.collections.PathCollection at 0x7fa0d35a7dd0>
```



```
#Perform Bivariate Analysis
```

```
df.corr()
```

	CreditScore	Gender	Age	Tenure	Balance \
CreditScore	1.000000	0.007888	-0.003965	0.000842	0.006268
Gender	0.007888	1.000000	0.022812	0.003739	0.069408
Age	-0.003965	0.022812	1.000000	-0.009997	0.028308
Tenure	0.000842	0.003739	-0.009997	1.000000	-0.012254
Balance	0.006268	0.069408	0.028308	-0.012254	1.000000
NumOfProducts	0.012238	0.003972	-0.030680	0.013444	-0.304180
HasCrCard	-0.005458	-0.008523	-0.011721	0.022583	-0.014858
IsActiveMember	0.025651	0.006724	0.085472	-0.028362	-0.010084
EstimatedSalary	-0.001384	-0.001369	-0.007201	0.007784	0.012797
Exited	-0.027094	0.035943	0.285323	-0.014001	0.118533

	NumOfProducts	HasCrCard	IsActiveMember	EstimatedSalary \
CreditScore	0.012238	-0.005458	0.025651	-
	0.001384			
Gender	0.003972	-0.008523	0.006724	-
	0.001369			
Age	-0.030680	-0.011721	0.085472	-
	0.007201			
Tenure	0.013444	0.022583	-0.028362	
	0.007784			
Balance	-0.304180	-0.014858	-0.010084	
	0.012797			
NumOfProducts	1.000000	0.003183	0.009612	
	0.014204			
HasCrCard	0.003183	1.000000	-0.011866	-
	0.009933			
IsActiveMember	0.009612	-0.011866	1.000000	-
	0.011421			
EstimatedSalary	0.014204	-0.009933	-0.011421	
	1.000000			
Exited	-0.047820	-0.007138	-0.156128	
	0.012097			

Exited

CreditScore	-0.027094	Gender
0.035943	Age	0.285323
Tenure	-0.014001	Balance
0.118533	NumOfProducts	-
0.047820		
HasCrCard	-0.007138	
IsActiveMember	-0.156128	
EstimatedSalary	0.012097	Exited
1.000000		

```
#Perform Bivariate Analysis
```

```
import statsmodels.api as sm
```

```
#define response
```

```
variable y =
```

```
df['CreditScore']
```

```
#define explanatory
```

```
variable x =
```

```
df[['EstimatedSalary']]
```

```
#add constant to predictor
```

```
variables x = sm.add_constant(x)
```

```
#fit linear regression model
```

```
model = sm.OLS(y, x).fit()
```

```
#view model summary
```

```
print(model.summary())
```

OLS

Regression Results

=====

=====

Dep. Variable: CreditScore R-squared: 0.000

Model: OLS Adj. R-squared:

-0.000

Method: Least Squares F-statistic:

0.01916

Date: Sat, 24 Sep 2022 Prob (F-statistic): 0.890

Time: 05:06:19 Log-Likelihood:

-59900.

No. Observations: 10000 AIC:

1.198e+05

Df Residuals: 9998 BIC:

1.198e+05

Df Model: 1

Covariance Type: nonrobust

```
=====
=====

            coef  std err          t      P>|t|
[0.025   0.975]
-----
-----

const          650.7617    1.940   335.407    0.000
646.958   654.565
EstimatedSalary -2.326e-06  1.68e-05   -0.138    0.890  -3.53e05   3.06e-05
=====
=====

Omnibus:          132.939  Durbin-Watson:
2.014
Prob(Omnibus):          0.000  Jarque-Bera (JB):
84.242
Skew:              -0.072  Prob(JB):
5.10e-19
Kurtosis:          2.574  Cond. No.
2.32e+05
=====
=====
```

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 2.32e+05. This might indicate that there are strong multicollinearity or other numerical problems.

/usr/local/lib/python3.7/dist-packages/statsmodels/tsa/tsatools.py:142:

FutureWarning: In a future version of pandas all arguments of concat except for the argument 'objs' will be keywordonly x = pd.concat(x[::order], 1)

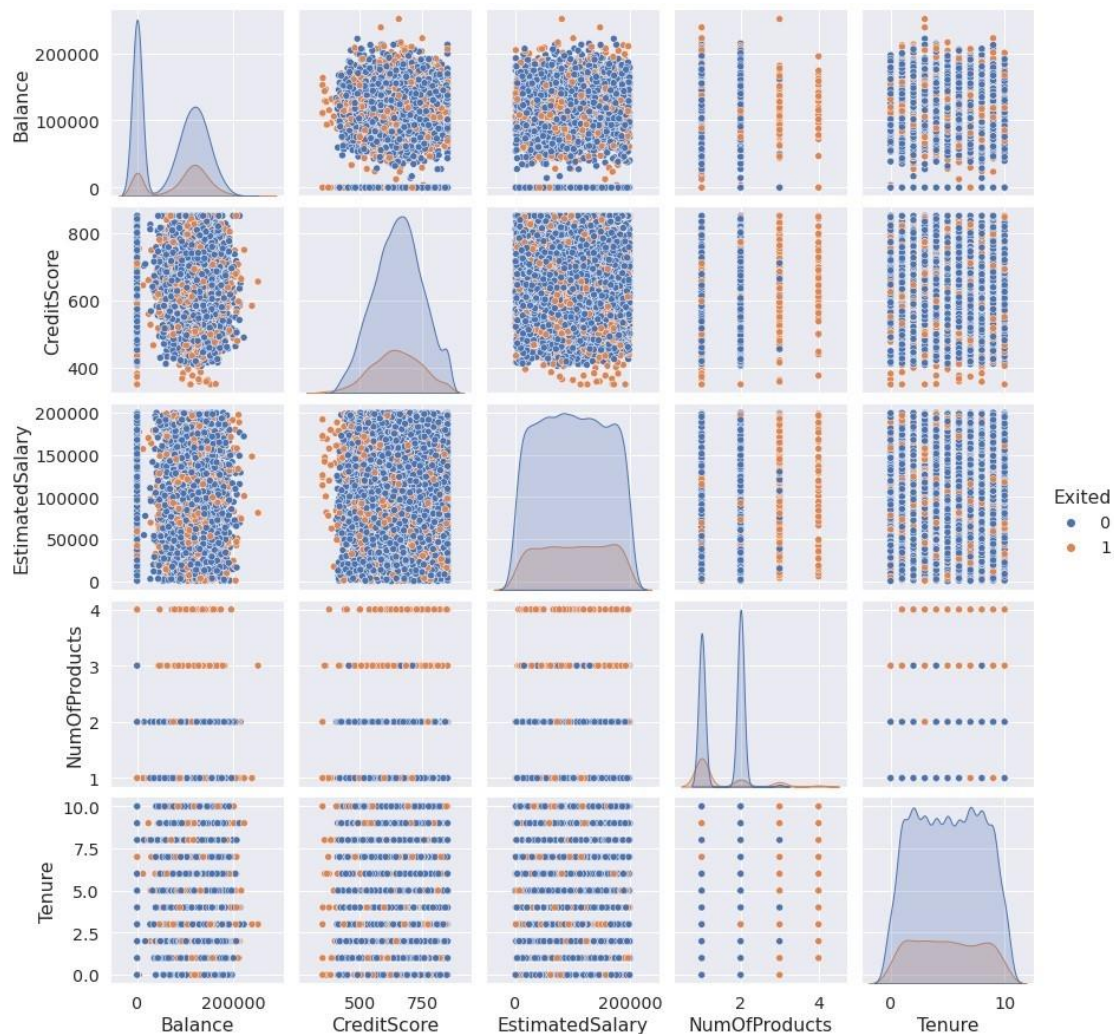
#Perform Multivariate Analysis

plt.figure(figsize=(4,4))

sns.pairplot(data=df[["Balance", "CreditScore", "EstimatedSalary", "NumOfProducts", "Tenure", "Exited"]], hue="Exited")

<seaborn.axisgrid.PairGrid at 0x7fa0b00a1b10>

<Figure size 288x288 with 0 Axes>



```
#Perform Descriptive  
Statistics df=pd.DataFrame(df)  
print(df.sum())
```

CreditScore 6505288 Geography
FranceSpainFranceFranceSpainSpainFranceGermany...

Gender FemaleFemaleFemaleFemaleFemaleMaleMaleFemaleMa... Age
389218 Tenure 50128 Balance 764858892.88
NumOfProducts 15302
HasCrCard 7055

IsActiveMember 5151 EstimatedSalary
1000902398.81 Exited 2037 dtype: object

```
#Perform Descriptive Statistics print("--  
--Sum Value----") print(df.sum(1))
```

```
print("-----") print("----Product Value----") print(df.prod())
```

```
print("-----")
```

```
----Sum Value----
```

```
0      102015.88
1      197002.44
2      274149.37
3      94567.63
4      205492.92    ... 9995  97088.64
9996    159633.38
9997    42840.58
9998    168784.83
9999    169159.57
```

```
Length: 10000, dtype: float64
```

```
-----Product Value-----
```

```
CreditScore    0.0
```

```
Age           0.0 Tenure           0.0
```

```
Balance        0.0 NumOfProducts
0.0
```

```
HasCrCard      0.0
```

```
IsActiveMember 0.0
```

```
EstimatedSalary inf Exited
```

```
0.0 dtype: float64
```

```
-----/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:3:
```

FutureWarning: Dropping of nuisance columns in DataFrame reductions (with 'numeric_only=None') is deprecated; in a future version this will raise TypeError. Select only valid columns before calling the reduction.

This is separate from the ipykernel package so we can avoid doing imports until

```
/usr/local/lib/python3.7/dist-packages/numpy/core/_methods.py:52: RuntimeWarning:
overflow encountered in reduce
```

```
return umr_prod(a, axis, dtype, out, keepdims, initial, where) /usr/local/lib/python3.7/dist-
packages/ipykernel_launcher.py:6:
```

FutureWarning: Dropping of nuisance columns in DataFrame reductions (with 'numeric_only=None') is deprecated; in a future version this will raise TypeError. Select only valid columns before calling the reduction.

```
#Perform Descriptive Statistics print("-----Mean
Value-----") print(df.mean())

print("-----") print("-----Median Value-----")
print(df.median())

print("-----") print("-----Mode Value-----")
print(df.mode())

print("-----")
```

-----Mean Value-----

```
CreditScore      650.528800
Age              38.921800 Tenure
5.012800 Balance      76485.889288
NumOfProducts    1.530200
HasCrCard        0.705500
IsActiveMember    0.515100 EstimatedSalary
100090.239881 Exited      0.203700 dtype:
float64
```

-----Median Value-----

```
CreditScore      652.000
Age              37.000 Tenure
5.000 Balance      97198.540
NumOfProducts    1.000
HasCrCard        1.000
IsActiveMember    1.000 EstimatedSalary
100193.915 Exited      0.000 dtype:
float64
```

-----Mode Value-----

```
CreditScore Geography Gender Age Tenure Balance
NumOfProducts \
0      850  France  Male  37    2    0.0      1
```

```
HasCrCard IsActiveMember EstimatedSalary Exited 0 1 1
24924.92 0
```

```
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:3:
```

FutureWarning: Dropping of nuisance columns in DataFrame reductions (with 'numeric_only=None') is deprecated; in a future version this will raise TypeError. Select only valid columns before calling the reduction.

This is separate from the ipykernel package so we can avoid doing imports until

```
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:6:
```

FutureWarning: Dropping of nuisance columns in DataFrame reductions (with 'numeric_only=None') is deprecated; in a future version this will raise TypeError. Select only valid columns before calling the reduction.

```
#Handling with missing Values
df.isnull()#Checking values are null
```

```
CreditScore Geography Gender Age Tenure Balance
NumOfProducts \
0          False  False  False  False  False  False
False
1          False  False  False  False  False  False
False
2          False  False  False  False  False  False
False
3          False  False  False  False  False  False
False
4          False  False  False  False  False  False
False
...      ...      ...      ...      ...      ...      ...
9995      False  False  False  False  False  False
False
9996      False  False  False  False  False  False
False
9997      False  False  False  False  False  False
False
9998      False  False  False  False  False  False  False
```

9999 False False False False False False False

HasCrCard IsActiveMember EstimatedSalary Exited

0	False	False	False	False
1	False	False	False	False
2	False	False	False	False
3	False	False	False	False
4	False	False	False	False

...

9995	False	False	False	False
9996	False	False	False	False
9997	False	False	False	False
9998	False	False	False	False
9999	False	False	False	False

[10000 rows x 11 columns]

#Handling with missing Values

df.notnull()#Checking values are not null

CreditScore Geography Gender Age Tenure Balance

NumOfProducts \

0	True	True	True	True	True	True
True						

1	True	True	True	True	True	True
True						

2	True	True	True	True	True	True
True						

3	True	True	True	True	True	True
True						

4	True	True	True	True	True	True
True						

...

9995	True	True	True	True	True	True
True						

9996	True	True	True	True	True	True
True						

9997	True	True	True	True	True	True
True						

9998 True True True True True True

True

9999 True True True True True True True

HasCrCard IsActiveMember EstimatedSalary Exited

0 True True True True

1 True True True True

2 True True True True

3 True True True True

4 True True True True

9995 True True True True

9996 True True True True

9997 True True True True

9998 True True True True

9999 True True True True


[10000 rows x 11 columns]

#Find outliers & replace the outliers

sns.boxplot(df['Balance'])

/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation. FutureWarning

<matplotlib.axes._subplots.AxesSubplot at 0x7fa0af6dcf90>



```
#Find outliers & replace the outliers
print(np.where(df['Balance']>100000))
```

```
(array([ 2,  4,  5, ..., 9987, 9993, 9999]),)
```

```
#Find outliers & replace the outliers
from scipy import stats
import numpy as np
```

```
z = np.abs(stats.zscore(df["EstimatedSalary"]))
print(z)
```

```
0      0.021886
1      0.216534
2      0.240687
3      0.108918
4      0.365276    ... 9995    0.066419
9996    0.027988
9997    1.008643
9998    0.125231
9999    1.076370
```

```
Name: EstimatedSalary, Length: 10000, dtype: float64
```

```
#Check for categorical columns & performs encoding
from sklearn.preprocessing import LabelEncoder
df['Gender'].unique()
```

```
array(['Female', 'Male'], dtype=object)
```

```
#Check for categorical columns & performs encoding
df['Gender'].value_counts()
```

```
Male    5457 Female
4543
```

```
Name: Gender, dtype: int64
```

```
#Check for categorical columns & performs encoding
encoding=LabelEncoder()
df["Gender"]=encoding.fit_transform(df.iloc[:,1].values) df
```

```
    CreditScore Geography Gender Age  Tenure  Balance
NumOfProducts \
0              619  France    0  42    2    0.00
1
1              608  Spain    2  41    1  83807.86
1
2              502  France    0  42    8 159660.80
3
3              699  France    0  39    1    0.00
2
4              850  Spain    2  43    2 125510.82
1
...      ...      ...      ...      ...      ...
9995              771  France    0  39    5    0.00
2
9996              516  France    0  35   10  57369.61
1
9997              709  France    0  36    7    0.00
1
9998              772  Germany    1  42    3  75075.31
2
9999              792  France    0  28    4 130142.79    1

    HasCrCard IsActiveMember EstimatedSalary  Exited
0           1           1    101348.88    1
1           0           1    112542.58    0
2           1           0    113931.57    1
```


3	0	0	93826.63	0
4	1	1	79084.10	0
9995	1	0	96270.64	0
9996	1	1	101699.77	0
9997	0	1	42085.58	1
9998	1	0	92888.52	1
9999	1	0	38190.78	0

[10000 rows x 11 columns]

#Check for categorical columns & performs encoding

#Split the data into Dependent & Independent Variables print("-----
----Dependent Variables-----")

X=df.iloc[:,1:4] print(X)

print("-----") print("-----Independent
Variables-----")

Y=df.iloc[:,4] print(Y)

print("-----")

-----Dependent Variables-----

	Age	Tenure	Balance	0	42
2	0.00				
1	41	1	83807.86		
2	42	8	159660.80		
3	39	1	0.00		
4	43	2	125510.82 9995 39 5 0.00
9996	35	10	57369.61		
9997	36	7	0.00		
9998	42	3	75075.31		
9999	28	4	130142.79		

[10000 rows x 3 columns]

-----Independent Variables-----

0	1
1	1
2	3
3	2
4	1 ..
9995	2
9996	1

```
9997    1
9998    2
9999    1
```

```
Name: NumOfProducts, Length: 10000, dtype: int64
```

```
-----
#Scale the independent Variables
```

```
from sklearn.preprocessing import StandardScaler object=
StandardScaler() # standardization scale =
object.fit_transform(df) print(scale)
```

```
[[-0.32622142  0.29351742 -1.04175968 ...  0.97024255  0.02188649
  1.97716468]
```

```
[-0.44003595  0.19816383 -1.38753759 ...  0.97024255  0.21653375
 -0.50577476]
```

```
[-1.53679418  0.29351742  1.03290776 ... -1.03067011  0.2406869  1.97716468] ...
[ 0.60498839 -0.27860412  0.68712986 ...  0.97024255 -1.00864308
  1.97716468]
```

```
[ 1.25683526  0.29351742 -0.69598177 ... -1.03067011 -0.12523071
  1.97716468]
```

```
[ 1.46377078 -1.04143285 -0.35020386 ... -1.03067011 -1.07636976 -0.50577476]]
```

```
#Split the data into training & testing from
sklearn.model_selection import train_test_split
```

```
#Split the data into training &
testing
```

```
x_train,x_test,y_train,y_test = train_test_split(x, y,
test_size=4,random_state=4) x_train
```

```
const EstimatedSalary 2558 1.0
137903.54
```

```
7642 1.0 121765.00
```

```
8912 1.0 109470.34
```

```
3319 1.0 2923.61
```

```
6852 1.0 7312.25 ... ...
```

```
... 456 1.0 7666.73
```

```
6017 1.0 9085.00
```

709 1.0 147794.63

8366 1.0 102515.42

1146 1.0 54776.64

[9996 rows x 2 columns]

```
#Split the data into training &
testing x_test
```

const EstimatedSalary 1603 1.0
23305.85

8713 1.0 41248.80

4561 1.0 143317.42

6600 1.0 174123.16

```
#Split the data into training &
testing y_train
```

2558 727

7642 811

8912 623

3319 430

6852 600

... 456 733

6017 487

709 686

8366 637

1146 614

Name: CreditScore, Length: 9996, dtype: int64

```
#Split the data into training &
testing y_test
```

1603 576

8713 786

4561 562

6600 505

Name: CreditScore, dtype: int64

